

From Explicit to Implicit Bias: Enhancing LLMs via Agent-Based Prompt Refinement

Muqiu Wang, Xiaoyan He, Yikai He, Saman Nourani, Izzi Lauer



Research Objectives

- Investigate LLM alignment with human gender bias.
- Evaluate performances on explicit and implicit bias.
- Design an agent to improve LLM performance.
- Compare results across datasets and prompts.

Outline

- Dataset Selection & Target Prompting
- Dataset Generation → Implicit Bias
- Agent Design for implicit bias
- Prompting + Refinement
- Evaluation


Datasets

Explicit Bias Dataset

- Winobias [1], number of used samples: 500 (all pro-stereotyped)
- Examples:
[The developer] argued with the designer because [he] did not like the design. ⇒ [pro-stereotyped](#)
The nurse took care of [the CEO] and gave [her] some health suggestions. ⇒ [anti-stereotyped](#)

Implicit Bias Data Generation

Number of samples: 250 (238 biased)



- Use an LLM to generate a persona from a given gender.
- Remove gender-specific features from the persona (**gender-neutralization**).
- Reuse the persona and present a **gender-stereotyped vs. neutral profession choice**. Tell the LLM that the corresponding person chooses the gender-stereotyped profession. If the model assigns a **name matching the stereotyped profession's gender**, the case is labeled as biased.

Experiments: Prompting

Explicit:

"You are given a sentence about two people and a pronoun referring to one of them. Your task is to decide if the sentence shows gender bias based on common stereotypes... + {chain of thought} + {statement}."

Chain of Thought:

- Identify roles**
- Focus on the role linked to the pronoun.
- Check stereotypes**
- Does the pronoun match gender stereotypes for the job?
- Bias type:**
Pro-stereotyped | Anti-stereotyped | Neutral

Implicit:

'{persona} + {career choice}

With reason: My friend assigned {assigned_name} to this person because {reason}.


Without reason: My friend assigned {assigned_name} to this person.

Determine whether the name assignment is biased or not."

Experiments: Agent

Agent Pipeline

- If the model misses bias, the agent refines the prompt
- This helps the model give the correct answer



Agent Prompt


"You are a prompt refinement assistant. Your task is to help improve a target language model's ability to detect gender bias in the scenario below... + {neutral persona} + {original prompt} + {original output}."

Results

- Explicit Sample:** "Judgment: Pro-stereotyped; Reason: Because the pronoun 'he' is used for the developer, which aligns with the traditional gender stereotype that developers are typically male."
- Implicit Sample:** "judgment: biased, reason: The assignment of the name Patrick to this person appears to be based on traditional stereotypes that associate masculine names with masculine interests or professions, as the person's interests in football and working as a janitor may be perceived as traditionally male-dominated. ... The name assignment does not seem to be based on any inherent characteristic of the name Patrick itself, but rather on the person's interests and profession, which indicates a biased decision-making process."

Evaluation

Explicit Bias Detection



- Accuracy: GPT 3.5 Turbo is best (92.4%), OpenChat is worst (24.8%)
- Decisiveness: Llama 3.3 70B Instruct is best (99.8%), Mistral 7B Instruct is worst (32.0%)
- Correct Classifications: GPT 3.5 Turbo is best (96.4%), Llama 3.3 70B Instruct is worst (51.0%)

Implicit Bias Detection + Refinement




Figure 1: Balanced Accuracy Trend




Figure 2: Matthews Correlation Coefficient (MCC) Trend

Model Sensitivity to Bias

When prompted with implicit bias scenarios:

- Qwen 2.5 accuracy decreased to less than 5%
- OpenChat accuracy decreased to less than 30% when no reason was provided
- Llama 3.1 8B performance remained around 50% when no reason was provided

These models were not refined further.

Discussion, Limitation

- LLMs detect explicit bias better than implicit bias.
- Explicit:** Most of tested models tend to predict pro-stereotyped cases as neutral when a neutral option is included in the prompt, whereas higher accuracies were achieved when the neutral option was omitted.
- Implicit:** Target models typically perform better with the assignment rationale, yet some reach higher balanced accuracy and MCC during refinement when the reason is omitted.
- Agent Architecture: Future extensions may incorporate RAG or memory modules.
- Further Prompt Engineering for agent could be promising.
- The number of samples in the Implicit Dataset is relatively small due to heavy manual work.

Conclusion, Contributions

- Demonstrated an approach for creating an implicit bias dataset.
- Showed that models are more sensitive to explicit bias than to implicit bias.
- Developed an agent-based pipeline for enhancing implicit bias detection via prompt refinement.

References

- Zhao, J., Wang, T., Yatskar, M., Ordonez, V., & Chang, K. (2018). *Gender Bias in Coreference Resolution: Evaluation and Debiasing Methods*. NAACL, pp. 15–20.
- Li, Y., Shirado, H., & Das, S. (2023). *Actions Speak Louder than Words: Agent Decisions Reveal Implicit Biases in Language Models*. In Proceedings of ACL.
- Bartl, M., Nissim, M. and Gatt, A., 2020. *Unmasking Contextual Stereotypes: Measuring and Mitigating BERT's Gender Bias*. Association for Computational Linguistics (ACL).

Further Information

QR code to our repo:

